

# Polynomial Computation on GPU


Kevin Mueller

March 15, 2016

# Outline

- Overview of Process
- Setup
- Homomorphisms (Mod Reduction)
- Interpolation
- Evaluation (Kernel)
- Problems encountered
- Extensions (Multivariate)
- Demo

# Overview of Process



# Initial problem Setup

Suppose we have two polynomials  $a(x) = 7x + 5$ ,  $b(x) = 2x - 3$  and we wish to multiply them such that

$$c(x) = a(x)b(x)$$

We first need to compute bounds for our modular design

- Compute upper bound  $M$  such that our set of moduli  $\prod m_i \geq 2M$
- For multiplication of two polynomials this is  $M = 2\|a\|_\infty\|b\|_\infty$

$$(7)(3)(2)(2) = 84$$

Therefore we can take  $m = 3, 5, 7$ . However for this problem we will use  $m = 5, 7$  for simplicity.

- Compute maximum degree for result polynomial

$$\deg(c) = \deg(a) + \deg(b)$$

# Homomorphisms

There are two types of homomorphisms we will be dealing with:

- Modular:  $\Phi_m : \mathbf{Z}[x_1, \dots, x_v] \rightarrow \mathbf{Z}_m[x_1, \dots, x_v]$
- In practice this simply means to take the mod on all the coefficients in the polynomial for some prime  $m_i$ . Such as

$$a_5(x) = 2x + 0 \quad b_5(x) = 2x + 2$$

- Evaluation:  $\Phi_{x_i-\alpha} : \mathbf{D}[x_1, \dots, x_v] \rightarrow \mathbf{D}[x_1, \dots, x_{i-1}, x_{i+1}, \dots, x_v]$
- This implies to evaluate each modular polynomial at each evaluation point. For a few examples:

$$a_5(0) = 0 \quad b_5(0) = 2$$

# Interpolation

We can solve the interpolation problem with the vandermonde matrix  $\mathbf{V}$  as

$$\mathbf{V}\mathbf{a} = \mathbf{b}$$

where  $n$  is the number of evaluation points,  $\mathbf{V}$  is the Vandermonde matrix with  $(i, j)$ -th entry  $\alpha_j^i(i, j = 0, 1, \dots, n)$ ,  $\mathbf{b}$  is the vector with the  $i$ -th entry  $b_i(i = 0, 1, \dots, n)$ , and  $\mathbf{a}$  is the vector of unknown coefficients  $a_i(i = 0, 1, \dots, n)$ .

$$\begin{pmatrix} c_5(0) = 0 \bmod 5 \\ c_5(1) = 3 \bmod 5 \\ c_5(2) = 4 \bmod 5 \end{pmatrix} \rightarrow c_5(x) = 4x^2 + 4x$$

$$\begin{pmatrix} c_7(0) = 6 \bmod 7 \\ c_7(1) = 2 \bmod 7 \\ c_7(2) = 5 \bmod 7 \end{pmatrix} \rightarrow c_7(x) = 3x + 6$$

# Chinese Remaindering

For a set of residues  $r_i \in \mathbb{Z} (1 \leq i \leq k)$  and a set of positive prime moduli  $m_i \in \mathbb{Z}_{m_i} (1 \leq i \leq k)$ , we can compute a unique  $u \in \mathbb{Z}_m$  in the mixed radix form as

$$u = v_1 + v_2(m_1) + v_3(m_1m_2) + \cdots + v_n \left( \prod_{i=0}^{k-1} m_i \right)$$

Which we can rewrite the 2nd coefficient as

$M_1 = 1, M_i = m_1m_2 \dots m_{i-1} (i = 2, \dots, k)$ . Satisfying the system of congruences gives us

$$v_1 = r_1$$

$$v_2 = (r_2 - v_1)\gamma_2 \mod m_2$$

$$v_i = ((r_i - v_1)\gamma_i - (\gamma_2 M_2 c_i \mod m_i) - \cdots - (v_{i-1} M_{i-1} \gamma_i \mod m_i))$$

# Chinese Remaindering continued

In practice we will precompute the coefficients as

$$\gamma_i = (m_1 m_2 \dots m_{i-1})^{-1} \pmod{m_i}$$

.

Finally we can recover the full integer  $r$  using the recursion

$$r = \gamma_1 + m_1(\gamma_2 + m_2(\gamma_3 + m_3(\dots)))$$

.



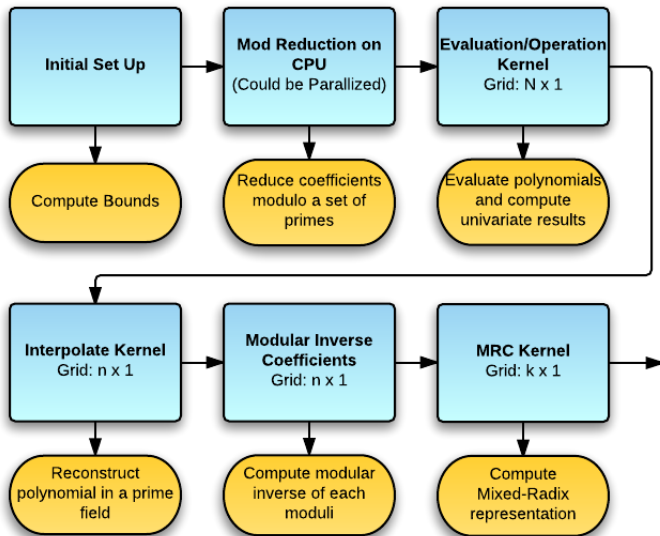


Figure: k: Number of Moduli, n: Number of Evaluation Points

# Evaluation Kernel Implementation

After performing all the modular homomorphisms necessary we want to maximize threads used to evaluate the polynomials in field  $\mathbf{Z}_{m_i}$ . In order to accomplish this we will need to vectorize all the coefficients and perform detailed bookkeeping to correctly index this coefficient vector  $X$ . We can then find the resulting vector  $Y \in \mathbb{Z}^N$  as

$$Y_i = a_j \alpha^\theta$$

Where  $\theta = (i \bmod d)$ ,  $j = \theta + d \cdot \text{quo}(i, n/c \cdot d)$  and  $\alpha = \text{quo}((i \bmod nd), d)$  when

- $l$  is the number of polynomials (i.e 2)
- $k = \text{len}(m)$
- $n = \text{deg}(c) + 1$
- $d = \max \{ \text{deg}(a(x)), \text{deg}(b(x)) \}$
- $N = l \cdot k \cdot n \cdot d$

# Problems Encountered

- Filling in zeros for coefficients with missing terms and handling string representations. Python libraries had poor support for this
- Python CUDA wrapper/compiler lacks some low level functionality.
- Coefficient swell in newton interpolation.
- Ill-conditioned Vandermonde matrix.
- Potential trade off between different polynomial representations.
- Sparsity required for operations in polynomial domain.
- Operation kernels for more complicated operations (division,GCD,resultant,etc).

# Extensions

There is a natural tendency to want to extend the above schemes into the multivariate domain. This seems somewhat straightforward since,

$$\mathbf{D}[x_3] \xrightarrow{\Phi_{x_2-\alpha}} \mathbf{D}[x_2, x_3] \xrightarrow{\Phi_{x_1-\alpha}} \mathbf{D}[x_1, x_2, x_3]$$

That is, each multivariate problem can be decomposed into a set of univariate interpolation problems where one evaluation homomorphism  $\Phi_{x_i-\alpha}$  is inverted at a time. However the interpolation becomes slightly less straight forward and sparsity becomes an even bigger issue.

A few other possible extensions could be:

- Perform the mod reductions on GPU
- Implement a parallel newton interpolation kernel
- Better support for mathematica

# References



Keith Geddes

Algorithms For Computer Algebra



Pavel Emeliyanenko

Harnessing the Power of GPUs for Problems in Real Algebraic Geometry.

*PhD thesis, Universitat des Saarlandes, 2012*